

DS203_Customer_Segmentation

November 13, 2021

```
[5]: import numpy as np
import pandas as pd
```

```
[6]: np.random.seed(42)
```

```
[7]: df = pd.read_csv('/content/Train.csv')
dataset = df.copy()
display(dataset)
```

	ID	Gender	Ever_Married	...	Family_Size	Var_1	Segmentation
0	462809	Male	No	...	4.0	Cat_4	D
1	462643	Female	Yes	...	3.0	Cat_4	A
2	466315	Female	Yes	...	1.0	Cat_6	B
3	461735	Male	Yes	...	2.0	Cat_6	B
4	462669	Female	Yes	...	6.0	Cat_6	A
...
8063	464018	Male	No	...	7.0	Cat_1	D
8064	464685	Male	No	...	4.0	Cat_4	D
8065	465406	Female	No	...	1.0	Cat_6	D
8066	467299	Female	No	...	4.0	Cat_6	B
8067	461879	Male	Yes	...	3.0	Cat_4	B

[8068 rows x 11 columns]

Performing basic data cleaning

```
[8]: dataset = dataset.drop(columns = ['ID'])
dataset = dataset[dataset['Profession'].notna()]
dataset = dataset.fillna(0)
dataset = dataset.loc[(dataset[['Graduated', 'Gender', 'Ever_Married', 'Age', 'Family_Size', 'Var_1']] != 0).all(axis=1)]

display(dataset)
```

	Gender	Ever_Married	Age	...	Family_Size	Var_1	Segmentation
0	Male	No	22	...	4.0	Cat_4	D

1	Female	Yes	38	...	3.0	Cat_4	A
2	Female	Yes	67	...	1.0	Cat_6	B
3	Male	Yes	67	...	2.0	Cat_6	B
4	Female	Yes	40	...	6.0	Cat_6	A
...
8062	Male	Yes	41	...	5.0	Cat_6	B
8064	Male	No	35	...	4.0	Cat_4	D
8065	Female	No	33	...	1.0	Cat_6	D
8066	Female	No	27	...	4.0	Cat_6	B
8067	Male	Yes	37	...	3.0	Cat_4	B

[7669 rows x 10 columns]

```
[9]: dataset_expanded = dataset.drop(columns = ['Segmentation'])
dataset_expanded = pd.get_dummies(dataset_expanded, sparse = True)
display(dataset_expanded)
```

	Age	Work_Experience	Family_Size	...	Var_1_Cat_5	Var_1_Cat_6	Var_1_Cat_7
0	22	1.0	4.0	...	0	0	0
1	38	0.0	3.0	...	0	0	0
2	67	1.0	1.0	...	0	1	0
3	67	0.0	2.0	...	0	1	0
4	40	0.0	6.0	...	0	1	0
...
8062	41	0.0	5.0	...	0	1	0
8064	35	3.0	4.0	...	0	0	0
8065	33	1.0	1.0	...	0	1	0
8066	27	1.0	4.0	...	0	1	0
8067	37	0.0	3.0	...	0	0	0

[7669 rows x 28 columns]

We can see that we have around 28 features now

```
[10]: from sklearn.preprocessing import StandardScaler
```

```
[11]: copy = dataset_expanded.copy()

scaler = StandardScaler()
scaler.fit(copy)
X_scale = scaler.transform(copy)
df_scale = pd.DataFrame(X_scale, columns=copy.columns)
display(df_scale)
```

	Age	Work_Experience	...	Var_1_Cat_6	Var_1_Cat_7
0	-1.288844	-0.412541	...	-1.40404	-0.160246
1	-0.330085	-0.712880	...	-1.40404	-0.160246

2	1.407665	-0.412541	...	0.71223	-0.160246
3	1.407665	-0.712880	...	0.71223	-0.160246
4	-0.210240	-0.712880	...	0.71223	-0.160246
...
7664	-0.150318	-0.712880	...	0.71223	-0.160246
7665	-0.509852	0.188138	...	-1.40404	-0.160246
7666	-0.629697	-0.412541	...	0.71223	-0.160246
7667	-0.989232	-0.412541	...	0.71223	-0.160246
7668	-0.390008	-0.712880	...	-1.40404	-0.160246

[7669 rows x 28 columns]

```
[12]: from sklearn.cluster import KMeans
      from sklearn.metrics import silhouette_score

[13]: df_scale2 = df_scale.copy()
      kmeans_scale = KMeans(n_clusters=4, n_init=100, max_iter=400, init='k-means++',
      →random_state=42).fit(df_scale2)
      print('KMeans Scaled Silhouette Score: {}'.format(silhouette_score(df_scale2,
      →kmeans_scale.labels_, metric='euclidean')))
      labels_scale = kmeans_scale.labels_
      clusters_scale = pd.concat([df_scale2, pd.DataFrame({'cluster_scaled':
      →labels_scale})], axis=1)
```

KMeans Scaled Silhouette Score: 0.1287897501946912

We can see that our Silhouette Score is extremely bad at 0.13 which suggests that there are too many features for KMeans

To solve this problem, we will be plotting a heatmap and see for what features the segmentation does not have a appreciable correlation so we can drop these features

```
[14]: dataset_heatmap = df.drop(columns = ['ID'])
      dataset_heatmap = dataset_heatmap[dataset_heatmap['Profession'].notna()]
      dataset_heatmap = dataset_heatmap.fillna(0)
      dataset_heatmap = dataset_heatmap.loc[(dataset_heatmap[['Graduated', 'Gender',
      →'Ever_Married', 'Var_1']] != 0).all(axis=1)]

      dataset_heatmap = pd.get_dummies(dataset_heatmap, sparse = True)
      display(dataset_heatmap)
```

	Age	Work_Experience	...	Segmentation_C	Segmentation_D
0	22	1.0	...	0	1
1	38	0.0	...	0	0
2	67	1.0	...	0	0
3	67	0.0	...	0	0
4	40	0.0	...	0	0
...


8062	41	0.0	...	0	0
8064	35	3.0	...	0	1
8065	33	1.0	...	0	1
8066	27	1.0	...	0	0
8067	37	0.0	...	0	0

[7669 rows x 32 columns]

```
[15]: import seaborn as sns
import matplotlib.pyplot as plt
plt.figure(figsize=(40, 30))

# plotting correlation heatmap
dataplot = sns.heatmap(dataset_heatmap.corr(), cmap="YlGnBu", annot=True)

# displaying heatmap
plt.show()
```



We can see that the segmentation does not really depend on factors like Gender and Work Experience and the Var_1 column. That means we can drop these columns for further analysis

```
[16]: dataset_cleaned = df.drop(columns = ['ID', 'Gender', 'Work_Experience', 'Var_1',
→ 'Segmentation'])
dataset_cleaned = dataset_cleaned[dataset_cleaned['Profession'].notna()]
dataset_cleaned = dataset_cleaned.loc[(dataset_cleaned[['Graduated',
→ 'Ever_Married']] != 0).all(axis=1)]
dataset_cleaned.Graduated.replace(('Yes', 'No'), (1, 0), inplace=True)
dataset_cleaned.Ever_Married.replace(('Yes', 'No'), (1, 0), inplace=True)
dataset_cleaned = dataset_cleaned.dropna()
display(dataset_cleaned)
```

	Ever_Married	Age	Graduated	Profession	Spending_Score	Family_Size
0	0.0	22	0.0	Healthcare	Low	4.0
1	1.0	38	1.0	Engineer	Average	3.0
2	1.0	67	1.0	Engineer	Low	1.0
3	1.0	67	1.0	Lawyer	High	2.0
4	1.0	40	1.0	Entertainment	High	6.0
...
8062	1.0	41	1.0	Artist	High	5.0
8064	0.0	35	0.0	Executive	Low	4.0
8065	0.0	33	1.0	Healthcare	Low	1.0
8066	0.0	27	1.0	Healthcare	Low	4.0
8067	1.0	37	1.0	Executive	Average	3.0

[7434 rows x 6 columns]

```
[17]: dataset_cleaned = pd.get_dummies(dataset_cleaned, sparse = True)
```

```
[18]: display(dataset_cleaned)
```

	Ever_Married	Age	...	Spending_Score_High	Spending_Score_Low
0	0.0	22	...	0	1
1	1.0	38	...	0	0
2	1.0	67	...	0	1
3	1.0	67	...	1	0
4	1.0	40	...	1	0
...
8062	1.0	41	...	1	0
8064	0.0	35	...	0	1
8065	0.0	33	...	0	1
8066	0.0	27	...	0	1
8067	1.0	37	...	0	0

[7434 rows x 16 columns]

We've successfully managed to reduce the dataset to 16 features from 28

```
[19]: copy2 = dataset_cleaned.copy()

scaler2 = StandardScaler()
scaler2.fit(copy2)
X_scale2 = scaler2.transform(copy2)
df_scale3 = pd.DataFrame(X_scale2, columns=copy2.columns)
display(df_scale3)
```

	Ever_Married	Age	...	Spending_Score_High	Spending_Score_Low
0	-1.205933	-1.295875	...	-0.425367	0.822321
1	0.829233	-0.332744	...	-0.425367	-1.216070
2	0.829233	1.412930	...	-0.425367	0.822321
3	0.829233	1.412930	...	2.350910	-1.216070
4	0.829233	-0.212353	...	2.350910	-1.216070
...
7429	0.829233	-0.152157	...	2.350910	-1.216070
7430	-1.205933	-0.513331	...	-0.425367	0.822321
7431	-1.205933	-0.633723	...	-0.425367	0.822321
7432	-1.205933	-0.994897	...	-0.425367	0.822321
7433	0.829233	-0.392940	...	-0.425367	-1.216070

[7434 rows x 16 columns]

```
[20]: df_scale4 = df_scale3.copy()
kmeans_scale = KMeans(n_clusters=4, n_init=100, max_iter=400, init='k-means++',
    random_state=42).fit(df_scale4)
print('KMeans Scaled Silhouette Score: {}'.format(silhouette_score(df_scale4,
    kmeans_scale.labels_, metric='euclidean')))
labels_scale = kmeans_scale.labels_
clusters_scale = pd.concat([df_scale4, pd.DataFrame({'cluster_scaled':
    labels_scale})], axis=1)
```

KMeans Scaled Silhouette Score: 0.23238024254207695

We can see a clear increase in our silhouette score to 0.23 from 0.13

```
[21]: display(clusters_scale)
```

	Ever_Married	Age	...	Spending_Score_Low	cluster_scaled
0	-1.205933	-1.295875	...	0.822321	3
1	0.829233	-0.332744	...	-1.216070	2
2	0.829233	1.412930	...	0.822321	1
3	0.829233	1.412930	...	-1.216070	0
4	0.829233	-0.212353	...	-1.216070	0
...
7429	0.829233	-0.152157	...	-1.216070	0

7430	-1.205933	-0.513331	...	0.822321	1
7431	-1.205933	-0.633723	...	0.822321	3
7432	-1.205933	-0.994897	...	0.822321	3
7433	0.829233	-0.392940	...	-1.216070	2

[7434 rows x 17 columns]

Try to perform elbow, other methods to find suitable clusters

1 PCA optional section

```
[22]: from sklearn.decomposition import PCA
pca = PCA(n_components = 0.95)
pca.fit(df_scale4)
reduced = pca.transform(df_scale4)
```

```
[27]: print(reduced.shape)
```

(7434, 12)

2 T-SNE

```
[28]: from sklearn.manifold import TSNE
import timeit
```

3 Two Dimensional Plot


```
[29]: start = timeit.default_timer()
tsne = TSNE(n_components=2, verbose=1, perplexity=80, n_iter=5000,
→learning_rate=200, random_state = 42)

tsne_scale_results = tsne.fit_transform(copy2)
end = timeit.default_timer()
tsne_df_scale = pd.DataFrame(tsne_scale_results, columns=['tsne1', 'tsne2'])
print('t-SNE done! Time elapsed: {} seconds'.format(end-start))
```

```
[t-SNE] Computing 241 nearest neighbors...
[t-SNE] Indexed 7434 samples in 0.018s...
[t-SNE] Computed neighbors for 7434 samples in 0.582s...
[t-SNE] Computed conditional probabilities for sample 1000 / 7434
[t-SNE] Computed conditional probabilities for sample 2000 / 7434
[t-SNE] Computed conditional probabilities for sample 3000 / 7434
[t-SNE] Computed conditional probabilities for sample 4000 / 7434
[t-SNE] Computed conditional probabilities for sample 5000 / 7434
```

```
[t-SNE] Computed conditional probabilities for sample 6000 / 7434
[t-SNE] Computed conditional probabilities for sample 7000 / 7434
[t-SNE] Computed conditional probabilities for sample 7434 / 7434
[t-SNE] Mean sigma: 0.967803
[t-SNE] KL divergence after 250 iterations with early exaggeration: 61.149101
[t-SNE] KL divergence after 5000 iterations: 0.692845
t-SNE done! Time elapsed: 368.51054587999994 seconds
```


```
[30]: plt.figure(figsize = (10,10))
plt.scatter(tsne_df_scale.iloc[:,0],tsne_df_scale.iloc[:,1],alpha=0.25,
           ↪facecolor='lightslategray')
plt.xlabel('tsne1')
plt.ylabel('tsne2')
plt.show()
```




```
[31]: kmeans_tsne_scale = KMeans(n_clusters=4, n_init=100, max_iter=800,
    →init='k-means++', random_state=42).fit(tsne_df_scale)
print('KMeans tSNE Scaled Silhouette Score: {}'.
    →format(silhouette_score(tsne_df_scale, kmeans_tsne_scale.labels_,
    →metric='euclidean'))
labels_tsne_scale = kmeans_tsne_scale.labels_
clusters_tsne_scale = pd.concat([tsne_df_scale, pd.DataFrame({'tsne_clusters':
    →labels_tsne_scale})], axis=1)
```

KMeans tSNE Scaled Silhouette Score: 0.5922470688819885

```
[32]: plt.figure(figsize = (10,10))
sns.scatterplot(x = clusters_tsne_scale.iloc[:,0],y = clusters_tsne_scale.iloc[:
    →,1],hue=labels_tsne_scale, palette='Set1', s=100, alpha=0.6).
    →set_title('Cluster Vis tSNE Scaled Data', fontsize=15)
plt.legend()
plt.show()
```



```
[41]: from sklearn.decomposition import PCA
```

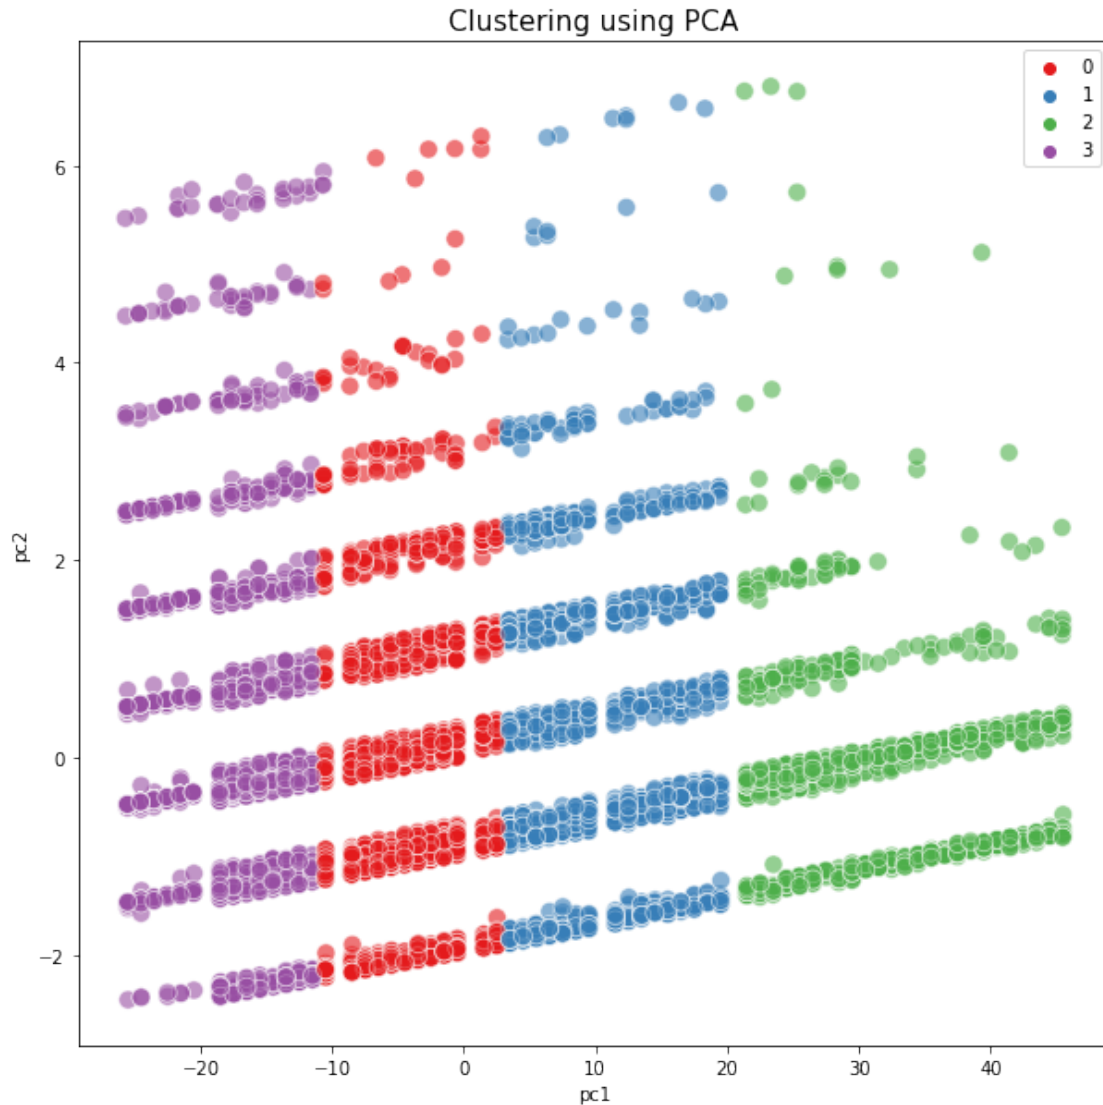
```
[46]: start = timeit.default_timer()
pca = PCA(n_components = 2, random_state = 42)
pca_scale_results = pca.fit_transform(copy2)
end = timeit.default_timer()
pca_df_scale = pd.DataFrame(pca_scale_results, columns=['pc1', 'pc2'])
print('t-SNE done! Time elapsed: {} seconds'.format(end-start))
print(pca.explained_variance_ratio_)
```

```
t-SNE done! Time elapsed: 0.03451279799992335 seconds
[0.98642641 0.00779754]
```

```
[43]: kmeans_pca_scale = KMeans(n_clusters=4, n_init=100, max_iter=800,
    →init='k-means++', random_state=42).fit(pca_df_scale)
print('KMeans PCA Scaled Silhouette Score: {}'.
    →format(silhouette_score(pca_df_scale, kmeans_pca_scale.labels_,
    →metric='euclidean'))
labels_pca_scale = kmeans_pca_scale.labels_
clusters_pca_scale = pd.concat([pca_df_scale, pd.DataFrame({'pca_clusters':
    →labels_pca_scale})], axis=1)
```

KMeans PCA Scaled Silhouette Score: 0.5286711782004786

```
[44]: plt.figure(figsize = (10,10))
sns.scatterplot(x = clusters_pca_scale.iloc[:,0],y = clusters_pca_scale.iloc[:
    →,1],hue=labels_pca_scale, palette='Set1', s=100, alpha=0.6).
    →set_title('Clustering using PCA', fontsize=15)
plt.legend()
plt.show()
```



4 3 Dimensional Plot

```
[33]: start = timeit.default_timer()
tsne = TSNE(n_components=3, verbose=1, perplexity=80, n_iter=5000,
→learning_rate=200, random_state = 42)

tsne_scale_results = tsne.fit_transform(copy2)
end = timeit.default_timer()
tsne_df_scale = pd.DataFrame(tsne_scale_results, columns=['tsne1', 'tsne2',
→'tsne3'])
print('t-SNE done! Time elapsed: {} seconds'.format(end-start))
```

```

[t-SNE] Computing 241 nearest neighbors...
[t-SNE] Indexed 7434 samples in 0.018s...
[t-SNE] Computed neighbors for 7434 samples in 0.618s...
[t-SNE] Computed conditional probabilities for sample 1000 / 7434
[t-SNE] Computed conditional probabilities for sample 2000 / 7434
[t-SNE] Computed conditional probabilities for sample 3000 / 7434
[t-SNE] Computed conditional probabilities for sample 4000 / 7434
[t-SNE] Computed conditional probabilities for sample 5000 / 7434
[t-SNE] Computed conditional probabilities for sample 6000 / 7434
[t-SNE] Computed conditional probabilities for sample 7000 / 7434
[t-SNE] Computed conditional probabilities for sample 7434 / 7434
[t-SNE] Mean sigma: 0.967803
[t-SNE] KL divergence after 250 iterations with early exaggeration: 60.717396
[t-SNE] KL divergence after 5000 iterations: 0.559766
t-SNE done! Time elapsed: 1293.319928696 seconds

```

```

[34]: kmeans_tsne_scale = KMeans(n_clusters=4, n_init=100, max_iter=800,
    →init='k-means++', random_state=42).fit(tsne_df_scale)
print('KMeans tSNE Scaled Silhouette Score: {}'.
    →format(silhouette_score(tsne_df_scale, kmeans_tsne_scale.labels_,
    →metric='euclidean'))))
labels_tsne_scale = kmeans_tsne_scale.labels_
clusters_tsne_scale = pd.concat([tsne_df_scale, pd.DataFrame({'tsne_clusters':
    →labels_tsne_scale})], axis=1)

```

KMeans tSNE Scaled Silhouette Score: 0.44695645570755005

```

[35]: import plotly.graph_objs as go

```

```

[36]: Scene = dict(xaxis = dict(title = 'tsne1'),yaxis = dict(title =
    →'tsne2'),zaxis = dict(title = 'tsne3'))
labels = labels_tsne_scale
trace = go.Scatter3d(x=clusters_tsne_scale.iloc[:,0], y=clusters_tsne_scale.
    →iloc[:,1], z=clusters_tsne_scale.iloc[:,2], mode='markers',marker=dict(color=
    →labels, colorscale='Viridis', size = 10, line = dict(color =
    →'yellow',width = 5)))
layout = go.Layout(margin=dict(l=0,r=0),scene = Scene, height = 1000,width =
    →1000)
data = [trace]
fig = go.Figure(data = data, layout = layout)
fig.show()

```

```

[48]: start = timeit.default_timer()
pca = PCA(n_components = 3, random_state = 42)
pca_scale_results = pca.fit_transform(copy2)
end = timeit.default_timer()
pca_df_scale = pd.DataFrame(pca_scale_results, columns=['pc1', 'pc2', 'pc3'])
print('PCA done! Time elapsed: {} seconds'.format(end-start))

```

```
print(pca.explained_variance_ratio_)
```

PCA done! Time elapsed: 0.04889983300017775 seconds
[0.98642641 0.00779754 0.00162805]

```
[49]: kmeans_pca_scale = KMeans(n_clusters=4, n_init=100, max_iter=800,
    →init='k-means++', random_state=42).fit(pca_df_scale)
print('KMeans PCA Scaled Silhouette Score: {}'.
    →format(silhouette_score(pca_df_scale, kmeans_pca_scale.labels_,
    →metric='euclidean'))
labels_pca_scale = kmeans_pca_scale.labels_
clusters_pca_scale = pd.concat([pca_df_scale, pd.DataFrame({'pca_clusters':
    →labels_pca_scale})], axis=1)
```

KMeans PCA Scaled Silhouette Score: 0.5225822486130006

```
[50]: Scene = dict(xaxis = dict(title = 'pc1'),yaxis = dict(title = 'pc2'),zaxis =
    →dict(title = 'pc3'))
labels = labels_pca_scale
trace = go.Scatter3d(x=clusters_pca_scale.iloc[:,0], y=clusters_pca_scale.iloc[:
    →,1], z=clusters_pca_scale.iloc[:,2], mode='markers',marker=dict(color =
    →labels, colorscale='Viridis', size = 10, line = dict(color = 'yellow',width
    →= 5)))
layout = go.Layout(margin=dict(l=0,r=0),scene = Scene, height = 1000,width =
    →1000)
data = [trace]
fig = go.Figure(data = data, layout = layout)
fig.show()
```

```
[ ]:
```